

Session #8: Device Drivers

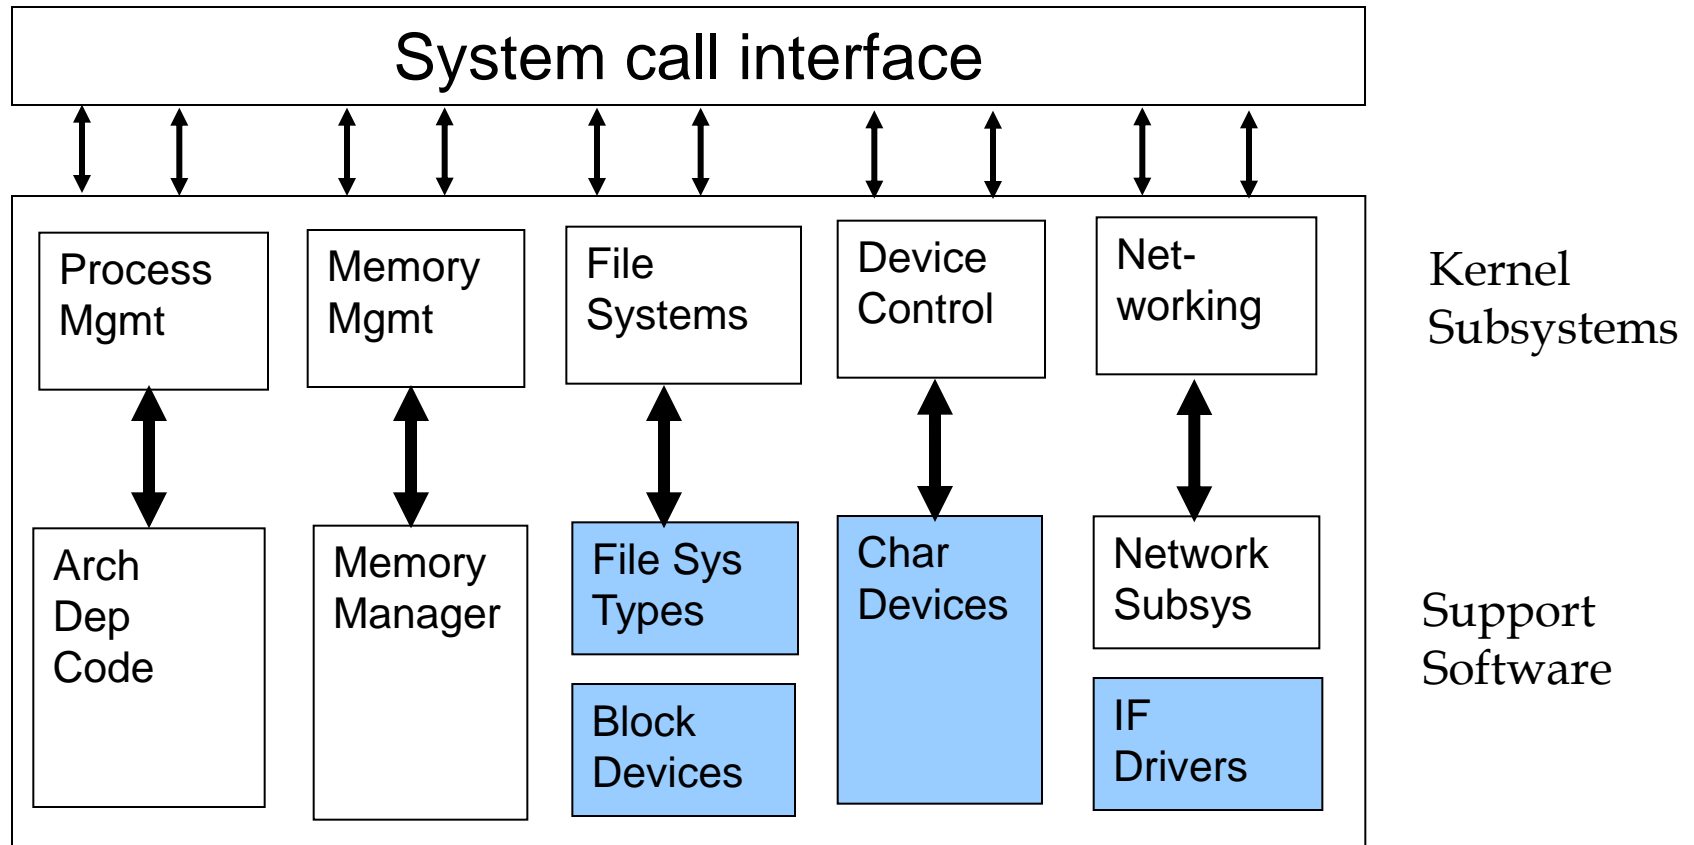
Programming at the Kernel Level



What's a "Device Driver" Anyway?

- A Set of Linkable Functions
- An Independently Loadable Program
- A Seperate Task or Set of Tasks
- A Well-defined driver API
- All of the Above??

Linux Architecture



Features that might be implemented as modules

Linux Device Drivers

- Devices treated like files
 - Everything in Linux is a file
 - Device files in /dev/...
- Device Classes
 - Character
 - Block
 - Pipe
 - Network

Low-level System Calls

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/ioctl.h>
```

```
int open (const char *path, int oflags);
size_t read (int file_des, void *data, size_t len);
size_t write (int file_des, const void *data, size_t len);
int close (int file_des);
int ioctl (int file_des, int cmd, ...);
```

Problems

- Kernel calls are inefficient
- Data isn't buffered

Standard I/O Library stdio

```
#include <stdio.h>
```

```
FILE *fopen (const char *filename, const char *mode);  
size_t fread (void *buff, size_t size, size_t n, FILE *stream);  
size_t fwrite (const void *buff, size_t size, size_t n, FILE *stream);  
int fclose (FILE *stream);  
getc, putc, and gets families  
Formatted I/O: printf, etc.
```

Problems

- Not good for device control (no ioctl function)
- Not deterministic

Installable Kernel Modules are...

- A way to “extend” the kernel
- Dynamically loaded and unloaded
- Executed at Privilege Level 0
- Useful for:
 - Device Drivers
 - Real-time tasks

Installable Kernel Modules

- Shell Commands
 - `insmod <module name> [<param>...]`
 - `rmmod <module name>`
- Required Functions
 - module initialization
 - module cleanup

Installing a Module

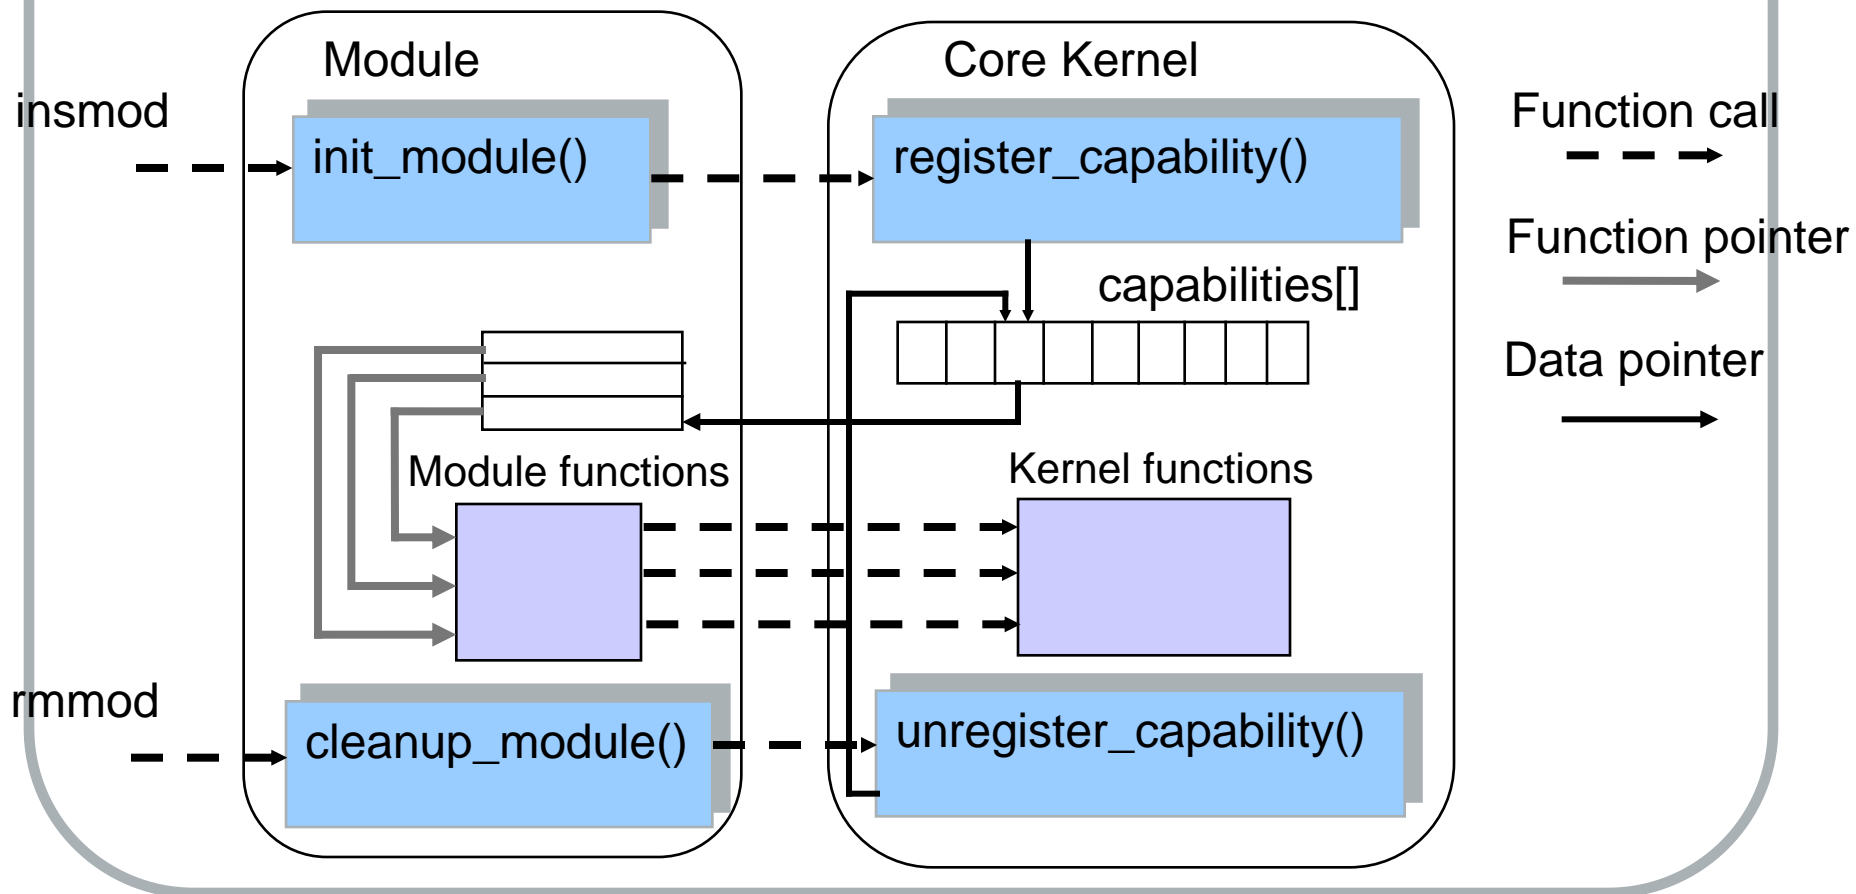
Install command

```
insmod my_module.ko my_int = 1 my_string = "Hello"
```

In my_module.c

```
int my_int;  
char *my_string;  
  
module_param (my_string, charp, S_IRUGO);  
module_param (my_int, int, S_IRUGO);  
  
static int __init my_init (void)  
{  
    register_capability();  
    return 0;  
}  
module_init (my_init);
```

Linking a Module to the Kernel



A Module Example

- `cd ~/drivers/hello`
- Look at `hello.c` and `Makefile`
- `make`
- `/sbin/lsmmod`
- In another terminal window, as root user
 - `tail -f /var/log/messages`
- `/sbin/insmod hello.ko my_string="your name" my_int=47`
 - What's wrong?
 - `MODULE_LICENSE()` is missing
- `/sbin/rmmod`

Kernel modules and the GPL

- GPL = GNU General Public License
- Applications in User Space need not be GPL
- Kernel is GPL
 - Any code built into the kernel is GPL
- Kernel modules need not be GPL
 - Some developers disagree

Allocating Device Numbers

`dev_t` - a 32-bit number in kernel 2.6

`MKDEV(int major, int minor)` - creates a `dev_t`

`MAJOR(dev_t dev)`

`MINOR(dev_t dev)`

`int register_chrdev_region (dev_t first, unsigned int count, char *name);`

Allocates `count` device numbers starting at `first`, if available

or

`int alloc_chrdev_region (dev_t *dev, unsigned int firstminor,
 unsigned int count, char *name,);`

Dynamically assigns a major device number for `count` devices starting at `firstminor`, usually zero.

Registering a Device

In `init_module()`

```
#include <linux/cdev.h>
```

```
struct cdev *my_cdev = cdev_alloc ( );  
my_cdev->ops = &my_fops;
```

or

```
void cdev_init (struct cdev *my_cdev, struct file_operations *my_fops);
```

followed by

```
my_cdev.owner = THIS_MODULE;  
int cdev_add (struct cdev *my_cdev, dev_t num, unsigned int count);
```

File Operations Table

```
struct file_operations {  
    struct module *owner;  
    loff_t (*llseek)(struct file *, loff_t, int);  
    ssize_t (*read)(struct file *, char __user *, size_t, loff_t *);  
    ssize_t (*aio_read)(struct kiocb *, char __user *, size_t, loff_t *);  
    ssize_t (*write)(struct file *, const char __user *, size_t, loff_t *);  
    ssize_t (*aio_write)(struct kiocb *, const char __user *, size_t, loff_t *);  
    int (*readdir)(struct file *, void *, filedir_t);  
    unsigned int (*poll)(struct file *, struct poll_table_struct *);  
    int (*ioctl)(struct file *, unsigned int, unsigned long);  
    int (*mmap)(struct file *, struct vm_area_struct *);  
    int (*open)(struct inode *, struct file *);  
    int (*flush)(struct file *);  
    int (*release)(struct inode *, struct file *);
```

File Operations Table 2

```
int (*fsync)(struct file *, struct dentry *, int);
int (*aio_fsync)(struct kiocb *, int);
int (*fasync)(int, struct file *, int);
int (*lock)(struct file *, int, struct file_lock *);
ssize_t (*readv)(struct file *, const struct iovec *, unsigned long, loff_t *);
ssize_t (*writev)(struct file *, const struct iovec *, unsigned long, loff_t *);
ssize_t (*sendfile)(struct file *, loff_t *, size_t, read_actor_t, void *);
ssize_t (*sendpage)(struct file *, struct page *, int, size_t, loff_t *, int);
unsigned long (*get_unmapped_area )(struct file *, unsigned long,
unsigned long, unsigned long, unsigned long);
int (*check_flags)(int);
int (*dir_notify )(struct file *, unsigned long);
}
```


struct file

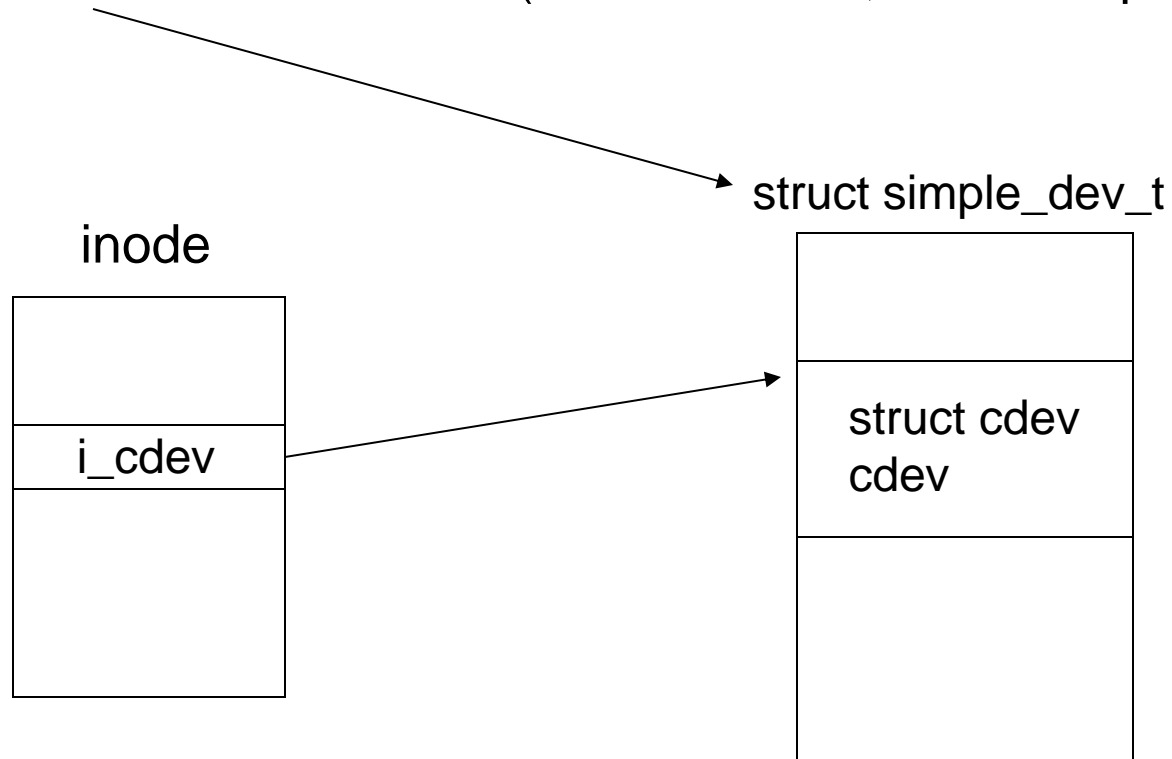
```
struct file {  
    struct list_head    f_list;  
    struct dentry       *f_dentry;  
    struct vsfmount     *f_vsfmount;  
    struct file_operations *f_op;  
    .  
    .  
    void *private_data;  
    .  
}
```

Our first driver

- `cd ~/drivers/simple_char`
- Look at `simple_char.c`
 - parameters, near top
 - `simple_init_module`, near bottom
 - `file_operations`
 - `simple_open`, back near top

container_of() macro

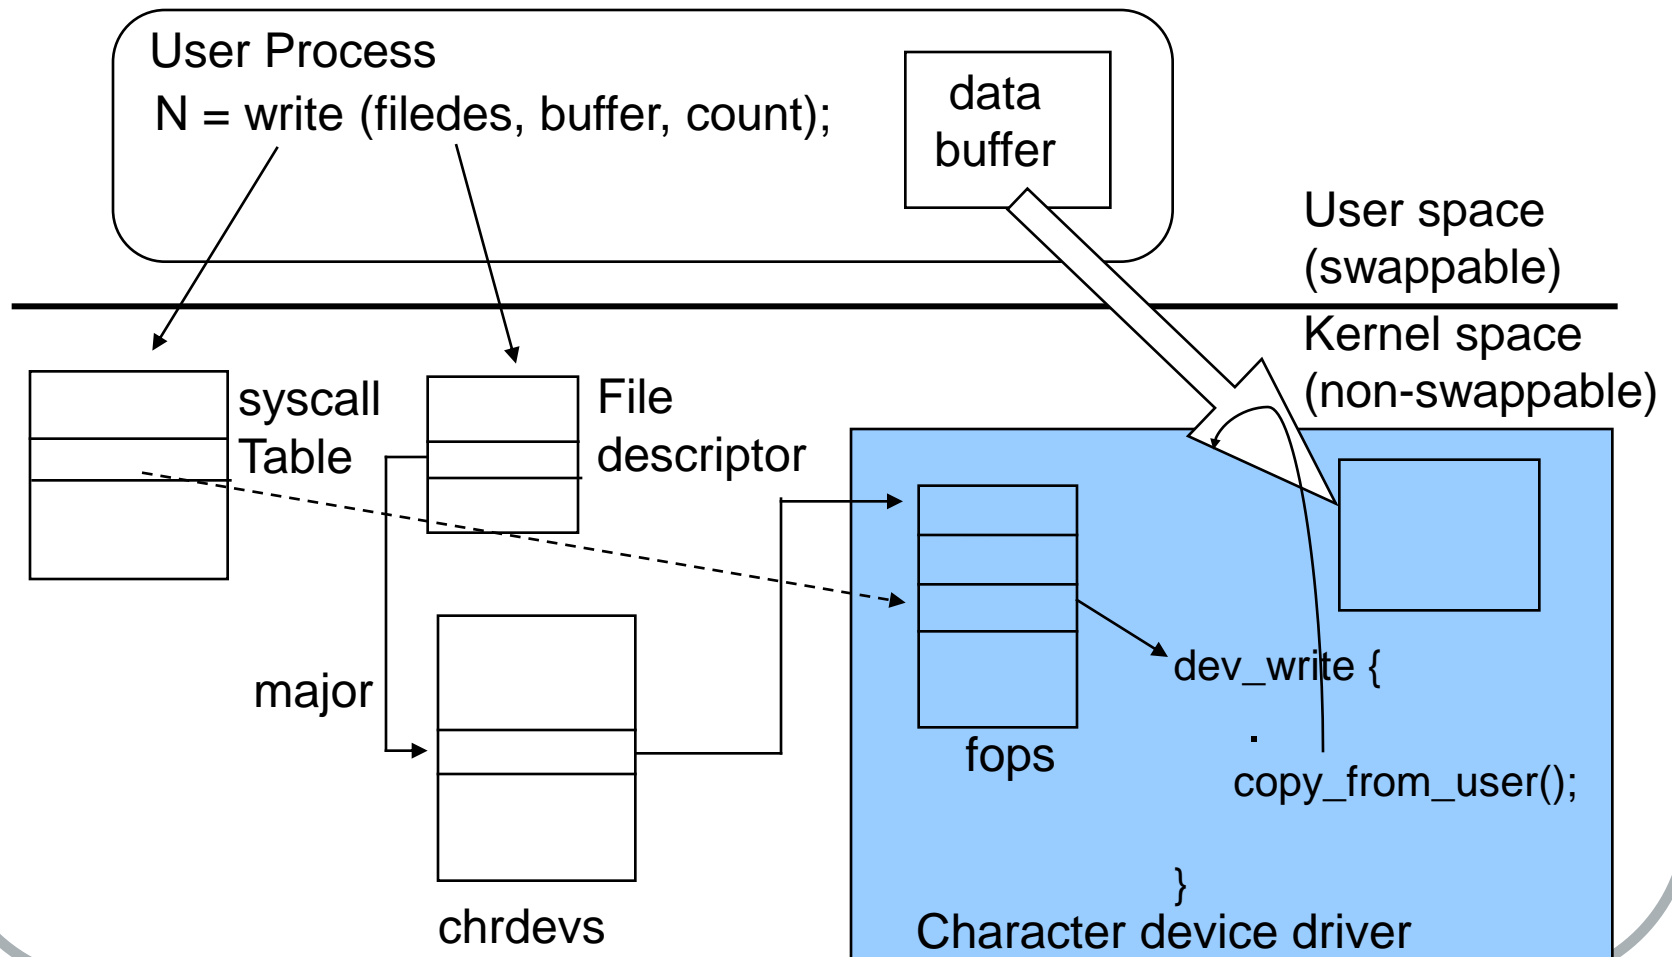
```
dev = container_of (inode->i_cdev, struct simple_dev_t, cdev);
```



Make and install simple_char

- make
- Look at `simple_load` script
 - Loads the driver module
 - Creates device nodes based on devices's major number
- `./simple_load`
- Try
 - `cp simple_char.c /dev/simple0`
 - `cat /dev/simple0`

Driver Processing--Write



Mutual exclusion--mutexes

- Only one process can access data at a time
- mutex guarantees exclusive access
 - `void mutex_init (struct mutex *);`
 - `void mutex_unlock (struct mutex *);`
 - `void mutex_lock (struct mutex *);`
 - `int mutex_lock_interruptible (struct mutex *);`

Registering a Device--the “old way”

In `init_module()`

```
int register_chrdev (unsigned int major, const char *name,  
                    struct file_operations *fops);
```

Registers driver at `chrdevs[major]`

or


```
int register_blkdev (unsigned int major, const char *name,  
                   struct file_operations *fops);
```

Registers driver at `blkdevs[major]`

printk()

```
driver_read(...)  
{  
    .  
    .  
    printk (KERN_DEBUG "In read function\n");  
    .  
    .  
}
```

loglevel



printk macro

In a header file

```
#ifdef DEBUG
#define PDEBUG (fmt, args, ...) printk (KERN_DEBUG fmt, ## args);
#else
#define PDEBUG (fmt, args, ...)    // nothing
#endif
```

```
driver_read(...)
{
    .
    .
    PDEBUG ( "In read function\n");
    .
    .
}
```

The /proc Filesystem

```
ls -l /proc
total 0
dr-xr-xr-x  3 root   root      0 Aug 25 15:23 1
dr-xr-xr-x  3 root   root      0 Aug 25 15:23 2
dr-xr-xr-x  3 root   root      0 Aug 25 15:23 3
dr-xr-xr-x  3 bin    root      0 Aug 25 15:23 303
dr-xr-xr-x  3 nobody nobody    0 Aug 25 15:23 416
dr-xr-xr-x  3 daemon daemon    0 Aug 25 15:23 434
dr-xr-xr-x  3 xfs    xfs      0 Aug 25 15:23 636
dr-xr-xr-x  4 root   root      0 Aug 25 15:23 bus
-r--r--r--  1 root   root      0 Aug 25 15:23 cmdline
-r--r--r--  1 root   root      0 Aug 25 15:23 cpuinfo
-r--r--r--  1 root   root      0 Aug 25 15:23 devices
-r--r--r--  1 root   root      0 Aug 25 15:23 filesystems
dr-xr-xr-x  2 root   root      0 Aug 25 15:23 fs
dr-xr-xr-x  4 root   root      0 Aug 25 15:23 ide
-r--r--r--  1 root   root      0 Aug 25 15:23 interrupts
-r--r--r--  1 root   root      0 Aug 25 15:23 ioports
```

Registering a /proc File

```
#include <linux/proc_fs.h>
```

/proc read function

```
int (*read_procmem) (struct file *filp, char __user *buf, size_t  
count, loff_t *pos);
```

Registering a /proc file

```
struct proc_dir_entry *proc_create (const char  
*name, mode_t mode, struct proc_dir_entry *parent,  
const struct file_operations *fops);
```

Removing a /proc file

```
void remove_proc_entry (const char *name, struct  
proc_dir_entry *parent);
```

ioctl--Another way to get info from driver

User Space

```
int ioctl (int fd, int cmd, ...);
```

Driver Method

```
int (*ioctl) (struct file *filp, unsigned int cmd,  
              unsigned long arg);
```

- Easier to code in driver than /proc file
- Requires User Space program to read or write
- Can leave in production code because nobody knows it's there

Choosing ioctl() commands

type - “Magic” number, 8 bits

 _IOCTL_TYPEBITS

number - Ordinal number for command, 8 bits

 __IOC_NRBITS

direction - From application’s viewpoint, bit mask

 _IOCTL_NONE

 _IOCTL_READ

 _IOCTL_WRITE

 _IOCTL_READ | _IOC_WRITE (both ways)

size - of item being transferred

Using pointer arguments

```
int access_ok (int type, const void *addr,  
              unsigned long size);
```

```
put_user (datum, ptr)  
__put_user (datum, ptr)
```

```
get_user (local, ptr)  
__get_user (local, ptr)
```

```
int capable (int capability);
```

Blocking I/O

- What if...
 - No data available (yet) on read?
 - Buffer full on write?
- Process must “sleep” until condition is satisfied
- Use wait queue (head)

Wait queue API

`#include <linux/wait.h>`

Create a wait queue

`DECLARE_WAIT_QUEUE_HEAD(name);`

or

`wait_queue_head_t my_queue;`
`init_waitqueue_head (&my_queue);`

Sleep on a queue

`wait_event (queue, condition)`
`wait_event_interruptible (queue, condition)`
`wait_event_timeout (queue, condition, timeout)`
`wait_event_interruptible_timeout (queue, condition, timeout)`

Wake up a process on a queue

`wake_up (&queue);`
`wake_up_interruptible (&queue);`

Putting a process to “sleep”

```
DECLARE_WAIT_QUEUE_HEAD(my_wait);
int flag = 0;
int buff_count;

ssize_t sleepy_write (struct file *file, char *buff, int count, loff_t *offset)
{
    copy_from_user (buffer, buff, buff_count = count);
    flag = 1;
    wake_up_interruptible (&my_wait);
    return count;
}

ssize_t sleepy_read (struct file *file, char *buff, int count, loff_t *offset)
{
    wait_event_interruptible (my_wait, flag != 0);
    copy_to_user (buffer, buff, buff_count);
    return buff_count;
}
```

Non-blocking operations

- `O_NONBLOCK` flag in `filp->f_flags`
- If set, functions return `-EAGAIN` if data transfer not possible
- Affects `open()`, `read()`, and `write()` file operations

The seq_file interface

```
void *seq_start (struct seq_file *sfile, loff_t *pos);  
void *seq_next (struct seq_file *sfile, void *v, loff_t *pos);  
void seq_stop (struct seq_file *sfile, void *v);  
int seq_show (struct seq_file *sfile, void *v);  
  
int seq_open (struct *file, struct seq_operations *seq_ops);
```

Review

- Device drivers
 - User space APIs
 - Device classes
- Loadable kernel modules
 - Modules and the GPL
- Basic character device driver
 - Debugging – `printk()`, `/proc` files, `ioctl`
- Blocking I/O